

EXNO : 1

Date

Install, configure and run Hadoop and HDFS

STEPS :

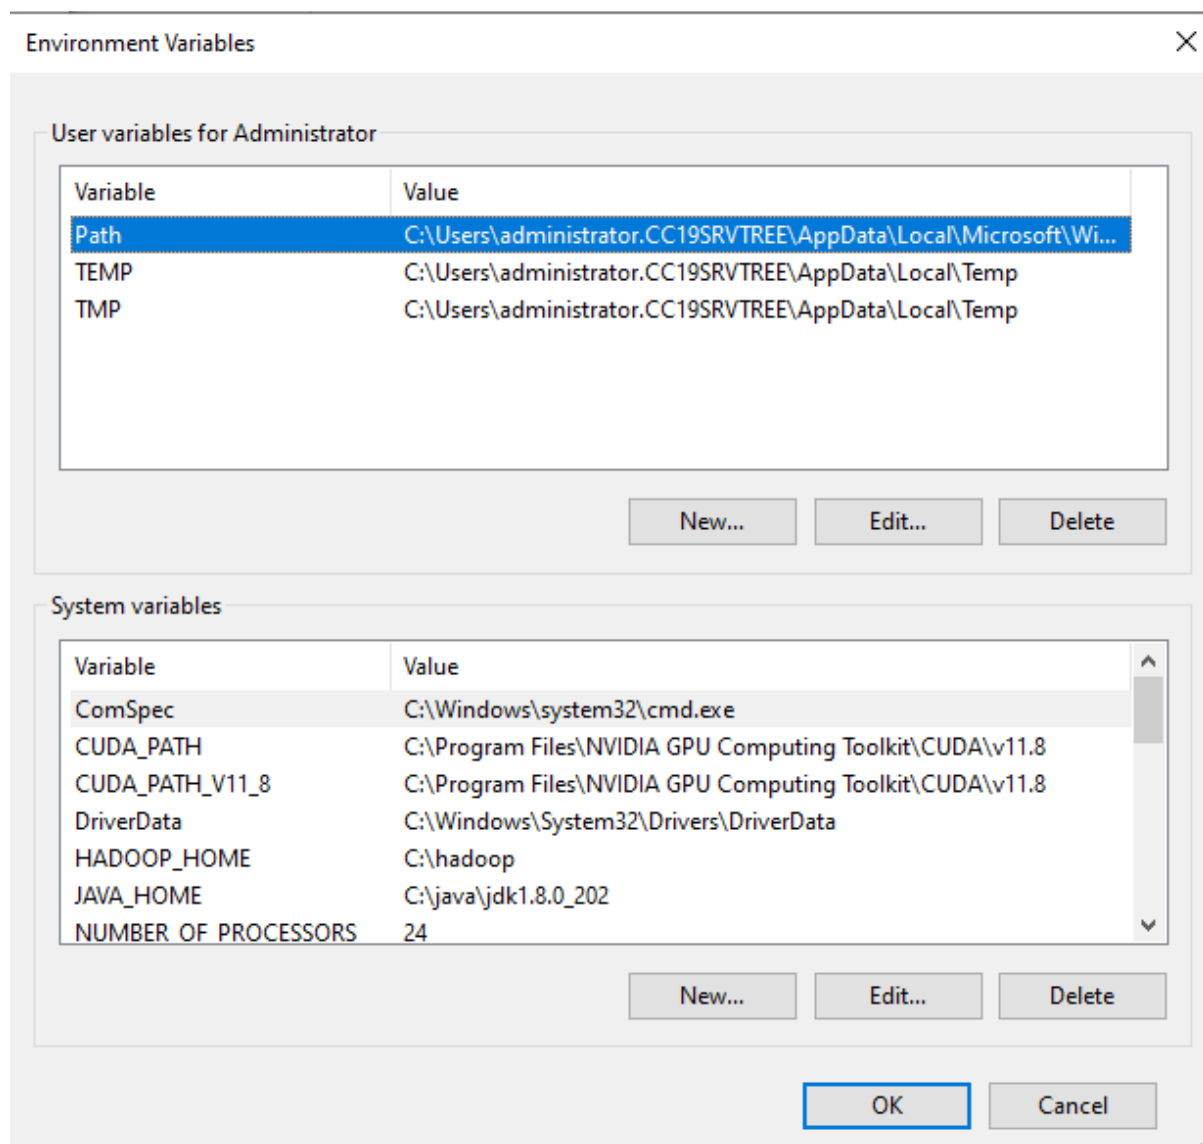
Step 1: Download Hadoop 3.2.4 get from

<https://hadoop.apache.org/release/3.2.4.html>

Step 2: Set the environment variable path for Hadoop

Go to Environment Variables and select it

User variable for administrator



Select new -> Create variable name is HADOOP_HOME = C:\hadoop\sbin

SYSTEMS VARIABLES

Select path and click edit -> new -> C:\hadoop\bin-> C:\hadoop\sbin

Step 3: Install java jdk 1.8.0 from

<https://www.oracle.com/in/java/technologies/javase/javase8-archive-downloads.html>

Set environment variable path for java like done for the Hadoop environmental variable setup

Go to Environment Variables and select it

User variable for administrator

Select new -> Create variable name is HOME_JAVA = C:\java\jdk1.8.0_202\bin

SYSTEMS VARIABLES

Select path and click edit -> new -> C:\java\jdk1.8.0_202\bin

Create a new folder named Java, and copy the JDK 1.8 folder from C:\Program Files\Java\jdk1.8.0_202 into the new Java folder

Step 4: Hadoop configuration

Set java path

Go to C:\hadoop\etc\hadoop -> hadoop-env -> @rem The java implementation to use. Required.

set JAVA_HOME=C:\java\jdk1.8.0_202

For core-site.xml

```
<property>
```

```
<name>fs.defaultFS</name>
```

```
<value>hdfs://localhost:9000</value>
```

```
</property>
```

For hdfs-site.xml or https-site.xml

```
<property> C:\hadoop\sbin
```

```
<name>dfs.replication</name>
```

```
<value>1</value>
```

```
</property><property>
```

```
<name>dfs.namenode.name.dir</name>
```

```
<value>C:\hadoop\data\namenode</value>
```

```
</property><property>
```

```
<name>dfs.datanode.data.dir</name>
```

```
<value>C:\hadoop\data\datanode</value>
```

```
</property>
```

For mapred-site.xml

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
```

For yarn-site.xml

```
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property><property>
  <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
```

Step:5 Creation of namenode and datanode

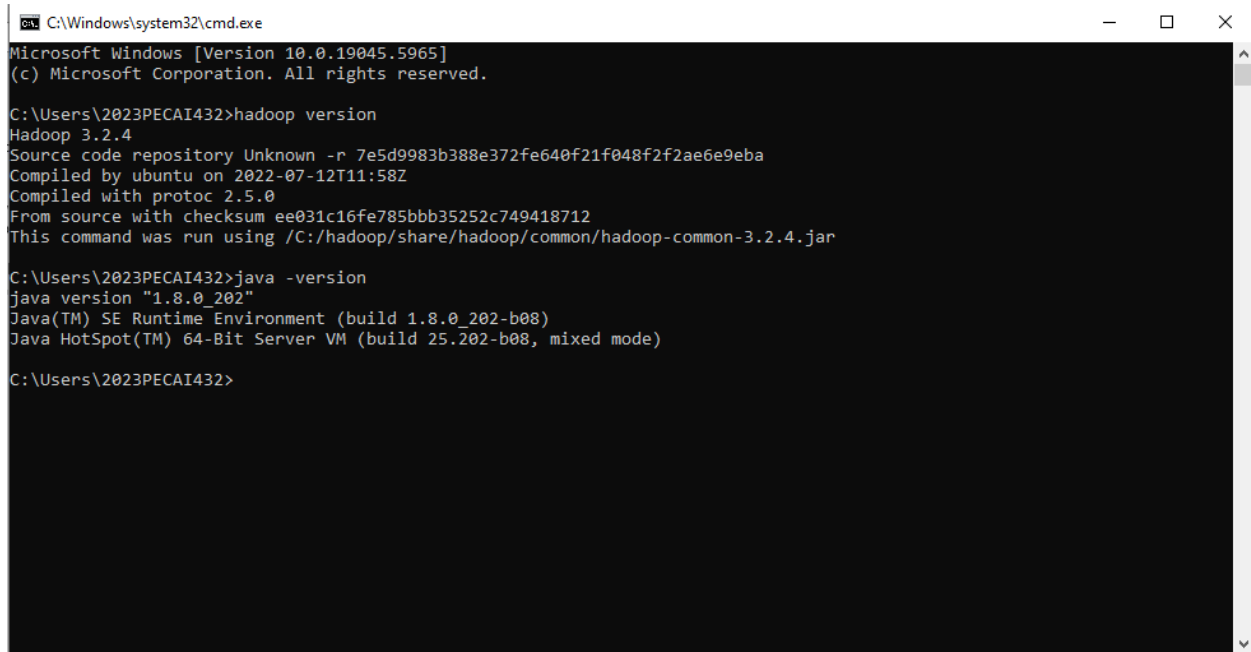
Create a new folder named data inside the Hadoop folder, and then create two subfolders inside the data folder named datanode and namenode

Step:6 Check whether Hadoop and Java are installed

Open Command Prompt and type it to check whether Hadoop and Java are installed

```
>>hadoop version
```

```
>>java -version
```



```
C:\Windows\system32\cmd.exe
Microsoft Windows [Version 10.0.19045.5965]
(c) Microsoft Corporation. All rights reserved.

C:\Users\2023PECAI432>hadoop version
Hadoop 3.2.4
Source code repository Unknown -r 7e5d9983b388e372fe640f21f048f2f2ae6e9eba
Compiled by ubuntu on 2022-07-12T11:58Z
Compiled with protoc 2.5.0
From source with checksum ee031c16fe785bbb35252c749418712
This command was run using /C:/hadoop/share/hadoop/common/hadoop-common-3.2.4.jar

C:\Users\2023PECAI432>java -version
java version "1.8.0_202"
Java(TM) SE Runtime Environment (build 1.8.0_202-b08)
Java HotSpot(TM) 64-Bit Server VM (build 25.202-b08, mixed mode)

C:\Users\2023PECAI432>
```

Step:7 Formatting Namenode and change directory from administrator to Hadoop sbin

Open Command Prompt and type

```
C:\Users\Administrator>hdfs namenode -format
```

Open Command Prompt as Administrator and navigate to the Hadoop sbin directory.

```
C:\Users\Administrator>cd C:\hadoop\sbin
```

Step:8 Now start all the Hadoop features

```
C:\hadoop\sbin>start-dfs.cmd
```

```
C:\hadoop\sbin>start-yarn.cmd
```

starting yarn daemons

```

Select Administrator: C:\Windows\system32\cmd.exe

0000000000000000 of size 408 bytes saved in 0 seconds .
2025-09-17 11:58:23,849 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2025-09-17 11:58:23,862 INFO namenode.FSNamesystem: Stopping services started for active state
2025-09-17 11:58:23,862 INFO namenode.FSNamesystem: Stopping services started for standby state
2025-09-17 11:58:23,866 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2025-09-17 11:58:23,867 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at sysdl16/172.16.37.49
*****/

C:\Users\Administrator>cd C:\hadoop\bin

C:\hadoop\bin>cd sbin
The system cannot find the path specified.

C:\hadoop\bin>cd C:\hadoop\sbin

C:\hadoop\sbin>start-dfs.cmd

C:\hadoop\sbin>start-yarn.cmd
starting yarn daemons

C:\hadoop\sbin>jps
10288 ResourceManager
9152 Jps
9696 NodeManager
9892 NameNode
11960 DataNode

C:\hadoop\sbin>

```

Step 9: Now open Microsoft Edge and go to localhost <http://localhost:9870> to check the Namenode interface in the Hadoop application

Overview 'localhost:9000' (active)

Started:	Wed Sep 17 12:05:48 +0530 2025
Version:	3.2.4, r7e5d9983b388e372fe640f21f048f2f2ae6e9eba
Compiled:	Tue Jul 12 17:28:00 +0530 2022 by ubuntu from branch-3.2.4
Cluster ID:	CID-9b021a68-519e-4a8e-b8f1-9232de7e6406
Block Pool ID:	BP-857010655-172.16.37.49-1758090503675

Summary

Security is off.
Safemode is off.
1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 74.02 MB of 292.5 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 51.58 MB of 53.65 MB Committed Non Heap Memory. Max Non Heap Memory is «unbounded».

Configured Capacity:	292.34 GB
Configured Remote Capacity:	0 B
DFS Used:	148 B (0%)
Non DFS Used:	225.57 GB

EXNO : 2

Date

Implement word count / frequency programs using MapReduce

Step 1: Formating the Namenode

Open Command Prompt as Administrator

```
C:\Users\Administrator>hdfs namenode -format
```

Step 2: Changing the directory from administrator to Hadoop sbin

Open Command Prompt as Administrator and navigate to the Hadoop sbin directory.

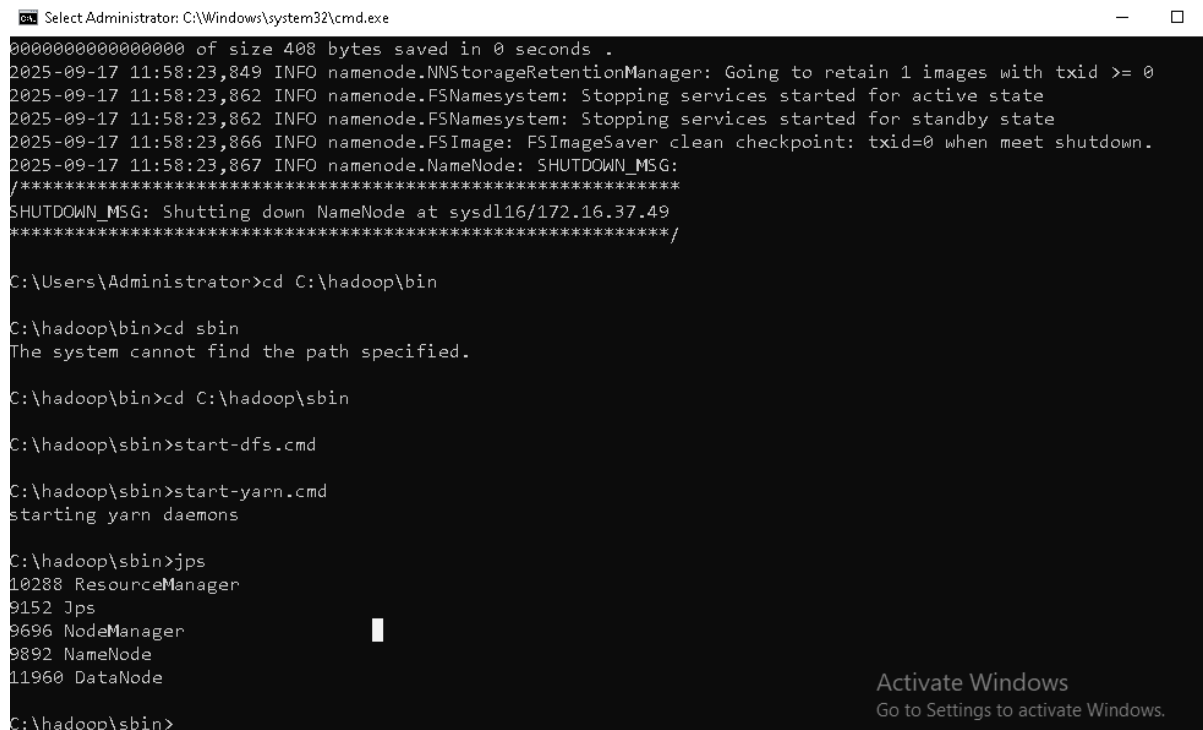
```
C:\Users\Administrator>cd C:\hadoop\sbin
```

Step 3: Start all Hadoop Features

```
C:\hadoop\sbin>start-dfs.cmd
```

```
C:\hadoop\sbin>start-yarn.cmd
```

starting yarn daemons



```
Select Administrator: C:\Windows\system32\cmd.exe

0000000000000000 of size 408 bytes saved in 0 seconds .
2025-09-17 11:58:23,849 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2025-09-17 11:58:23,862 INFO namenode.FSNamesystem: Stopping services started for active state
2025-09-17 11:58:23,862 INFO namenode.FSNamesystem: Stopping services started for standby state
2025-09-17 11:58:23,866 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2025-09-17 11:58:23,867 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at sysdl16/172.16.37.49
*****/

C:\Users\Administrator>cd C:\hadoop\bin

C:\hadoop\bin>cd sbin
The system cannot find the path specified.

C:\hadoop\bin>cd C:\hadoop\sbin

C:\hadoop\sbin>start-dfs.cmd

C:\hadoop\sbin>start-yarn.cmd
starting yarn daemons

C:\hadoop\sbin>jps
10288 ResourceManager
9152 Jps
9696 NodeManager
9892 NameNode
11960 DataNode

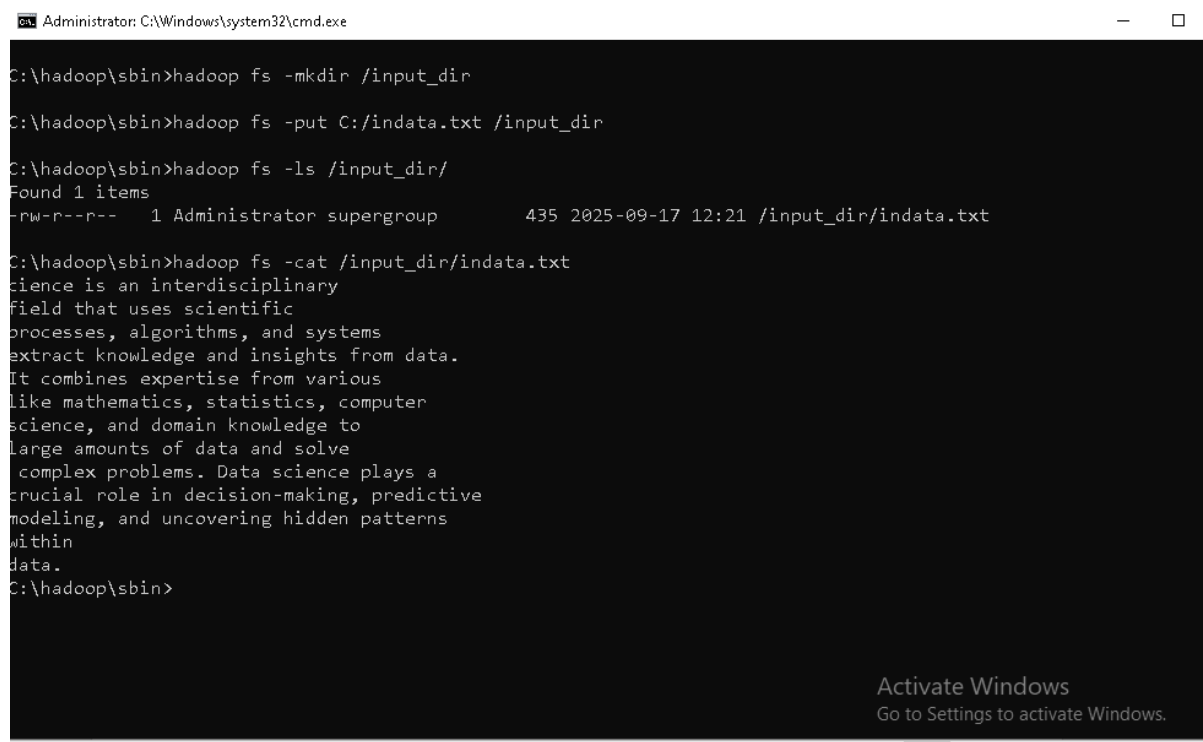
C:\hadoop\sbin>
```

Step 4: Create a indata.txt in the input directory

Open command prompt and type the following commands to create the indata.txt file in the directory named input

Use this directory for input file C:\hadoop\input_dir\indata.txt

```
>> hadoop fs -mkdir /input_dir
>> hadoop fs -put C:/indata.txt /input_dir
>> hadoop fs -ls /input_dir/
>> hadoop fs -cat /input_dir/indata.txt
```



```
Administrator: C:\Windows\system32\cmd.exe

C:\hadoop\sbin>hadoop fs -mkdir /input_dir

C:\hadoop\sbin>hadoop fs -put C:/indata.txt /input_dir

C:\hadoop\sbin>hadoop fs -ls /input_dir/
Found 1 items
-rw-r--r--  1 Administrator supergroup      435 2025-09-17 12:21 /input_dir/indata.txt

C:\hadoop\sbin>hadoop fs -cat /input_dir/indata.txt
science is an interdisciplinary
field that uses scientific
processes, algorithms, and systems
extract knowledge and insights from data.
It combines expertise from various
like mathematics, statistics, computer
science, and domain knowledge to
large amounts of data and solve
complex problems. Data science plays a
crucial role in decision-making, predictive
modeling, and uncovering hidden patterns
within
data.
C:\hadoop\sbin>
```

Activate Windows
Go to Settings to activate Windows.

Step 5: Save the output in the output directory

Hadoop jar C:/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.4.jar wordcount /input_dir /output_dir

Use this directory to get the wordcount from the input file in the output directory

```
Administrator: C:\Windows\system32\cmd.exe
data.
C:\hadoop\sbin>hadoop jar C:/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.4.jar wordcount /input_dir /output_dir
2025-09-17 12:53:21,235 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
2025-09-17 12:53:21,582 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/Administrator/.staging/job_1758091042685_0001
2025-09-17 12:53:21,694 INFO input.FileInputFormat: Total input files to process : 1
2025-09-17 12:53:21,743 INFO mapreduce.JobSubmitter: number of splits:1
2025-09-17 12:53:21,799 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1758091042685_0001
2025-09-17 12:53:21,800 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-09-17 12:53:21,922 INFO conf.Configuration: resource-types.xml not found
2025-09-17 12:53:21,922 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2025-09-17 12:53:22,276 INFO impl.YarnClientImpl: Submitted application application_1758091042685_0001
2025-09-17 12:53:22,322 INFO mapreduce.Job: The url to track the job: http://sysd116:8088/proxy/application_1758091042685_0001/
2025-09-17 12:53:22,322 INFO mapreduce.Job: Running job: job_1758091042685_0001
2025-09-17 12:53:28,424 INFO mapreduce.Job: Job job_1758091042685_0001 running in uber mode : false
2025-09-17 12:53:28,427 INFO mapreduce.Job: map 0% reduce 0%
2025-09-17 12:53:31,474 INFO mapreduce.Job: map 100% reduce 0%
2025-09-17 12:53:35,530 INFO mapreduce.Job: map 100% reduce 100%
2025-09-17 12:53:35,549 INFO mapreduce.Job: Job job_1758091042685_0001 completed successfully
2025-09-17 12:53:35,609 INFO mapreduce.Job: Counters: 54
File System Counters
  FILE: Number of bytes read=674
  FILE: Number of bytes written=479611
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=542
  HDFS: Number of bytes written=472
```

localhost:9870/explorer.html#/

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

Browse Directory

/

Go

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	Administrator	supergroup	0 B	Sep 17 12:21	0	0 B	input_dir
drwxr-xr-x	Administrator	supergroup	0 B	Sep 17 12:53	0	0 B	output_dir
drwx-----	Administrator	supergroup	0 B	Sep 17 12:53	0	0 B	tmp

Showing 1 to 3 of 3 entries

Previous 1 Next

Hadoop, 2022.

Activate Windows
Go to Settings to activate Windows.

OUTPUT:

hadoop fs -cat /output_dir/*



```
Administrator: C:\Windows\system32\cmd.exe
Bytes Written=472
C:\hadoop\sbin>hadoop fs -cat /output_dir/*
Data      1
It        1
a         1
algorithms, 1
amounts 1
an        1
and       5
cience 1
combines 1
complex 1
computer 1
crucial 1
data      1
data.    2
decision-making, 1
domain 1
expertise 1
extract 1
field 1
from 2
hidden 1
in 1
insights 1
interdisciplinary 1
is 1
knowledge 2
large 1
```

EXNO : 3

Date

Implement an MR program that processes a weather dataset

Step 1: Download netbeans 8.2 from

<https://netbeans-ide.informer.com/download/>

Step 2: Install the netbeans 8.2

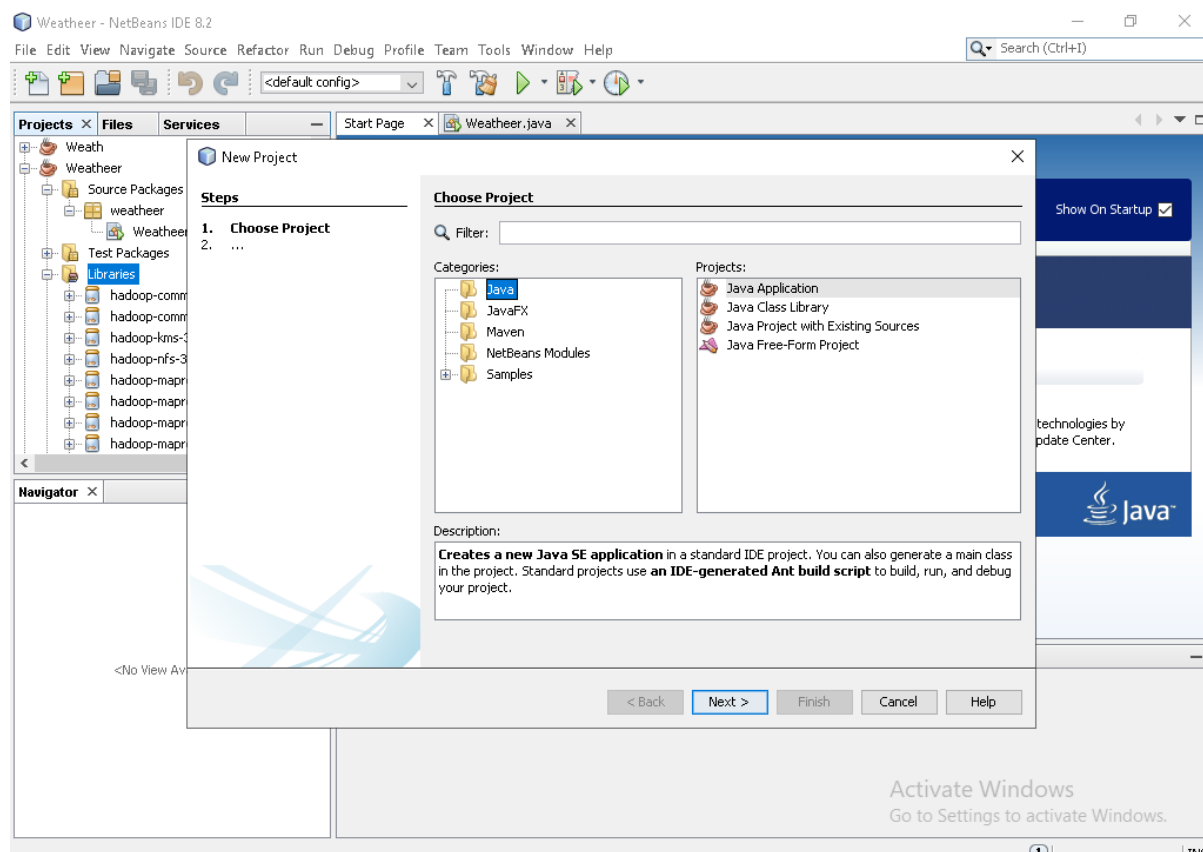
Install the downloaded netbeans 8.2 application in the system.

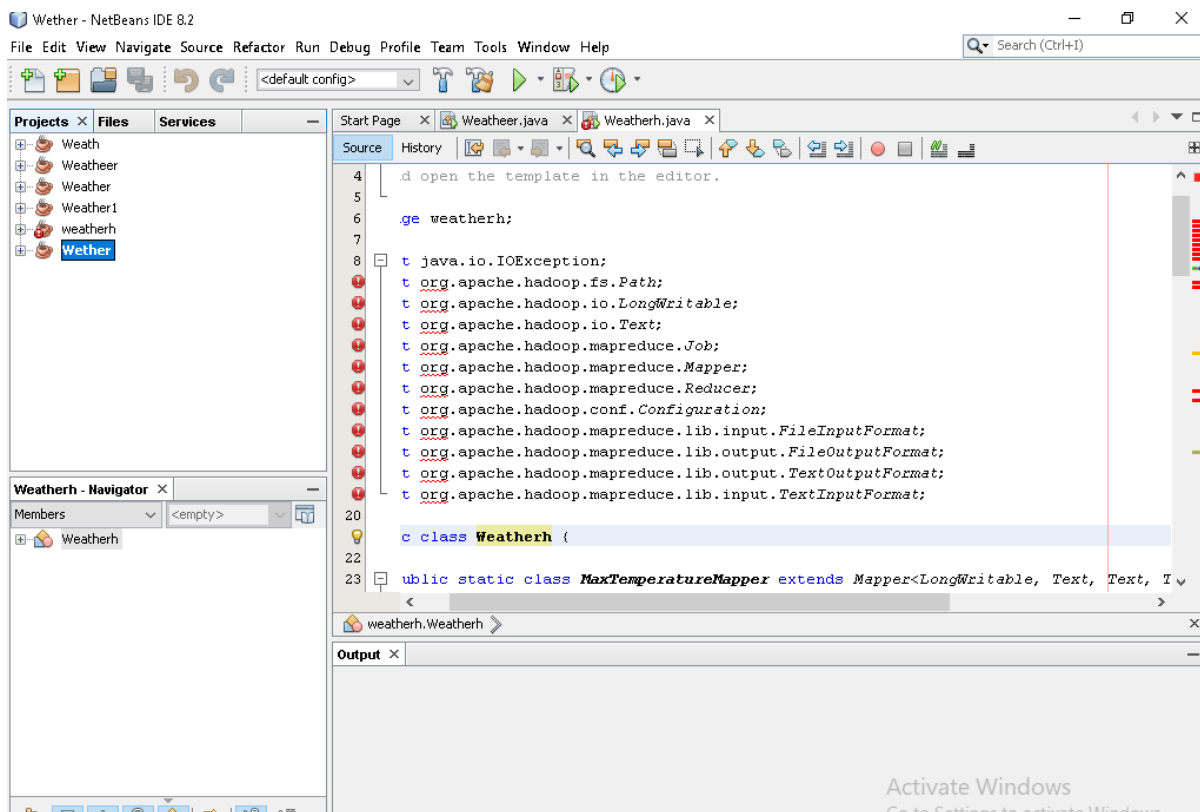
Step 3: Creating s new Project

Follow the steps below to create a new project in netbeans

File -> select new project -> categories -> java and project -> java application click on next

Enter your project name





#PROGRAM:

```
package weatheer;

import java.io.IOException;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.LongWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;

public class Weatheer {
```

```
public static class MaxTemperatureMapper extends Mapper<LongWritable, Text, Text, Text> {

    @Override

    public void map(LongWritable key, Text value, Context context) throws IOException,
    InterruptedException {

        String line = value.toString();

        if (line.trim().isEmpty()) return; // Skip empty lines

        try {

            String date = line.substring(6, 14);

            float tempMax = Float.parseFloat(line.substring(39, 45).trim());

            float tempMin = Float.parseFloat(line.substring(47, 53).trim());

            if (tempMax > 35.0) {

                context.write(new Text("Hot Day " + date), new Text(String.valueOf(tempMax)));

            }

            if (tempMin < 10.0) {

                context.write(new Text("Cold Day " + date), new Text(String.valueOf(tempMin)));

            }

        } catch (Exception e) {

            // Handle lines with invalid format or parsing errors

            // You can log this if needed

        }

    }

}

public static class MaxTemperatureReducer extends Reducer<Text, Text, Text, Text> {
```

```
@Override

public void reduce(Text key, Iterable<Text> values, Context context) throws IOException,
InterruptedException {

    // Simply write the first temperature found

    context.write(key, values.iterator().next());

}

}

public static void main(String[] args) throws Exception {

    if (args.length < 2) {

        System.err.println("Usage: Weather <input path> <output path>");

        System.exit(-1);

    }

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Weather Example");

    job.setJarByClass(Weather.class);

    job.setMapperClass(MaxTemperatureMapper.class);

    job.setReducerClass(MaxTemperatureReducer.class);

    job.setMapOutputKeyClass(Text.class);

    job.setMapOutputValueClass(Text.class);

    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(Text.class);

    job.setInputFormatClass(TextInputFormat.class);

    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));

    Path outputPath = new Path(args[1]);
```

```
// Delete output path if exists

outputPath.getFileSystem(conf).delete(outputPath, true);

FileOutputFormat.setOutputPath(job, outputPath);

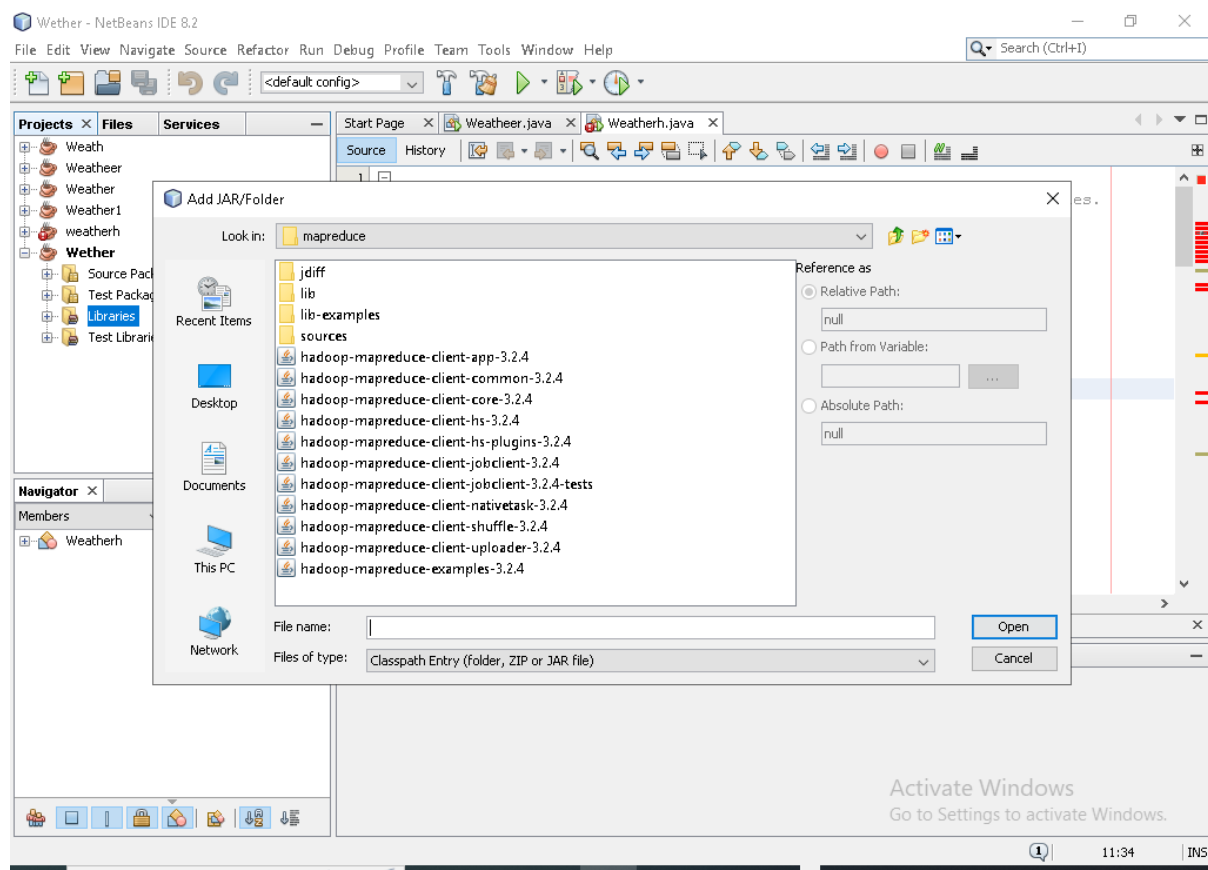
System.exit(job.waitForCompletion(true) ? 0 : 1);

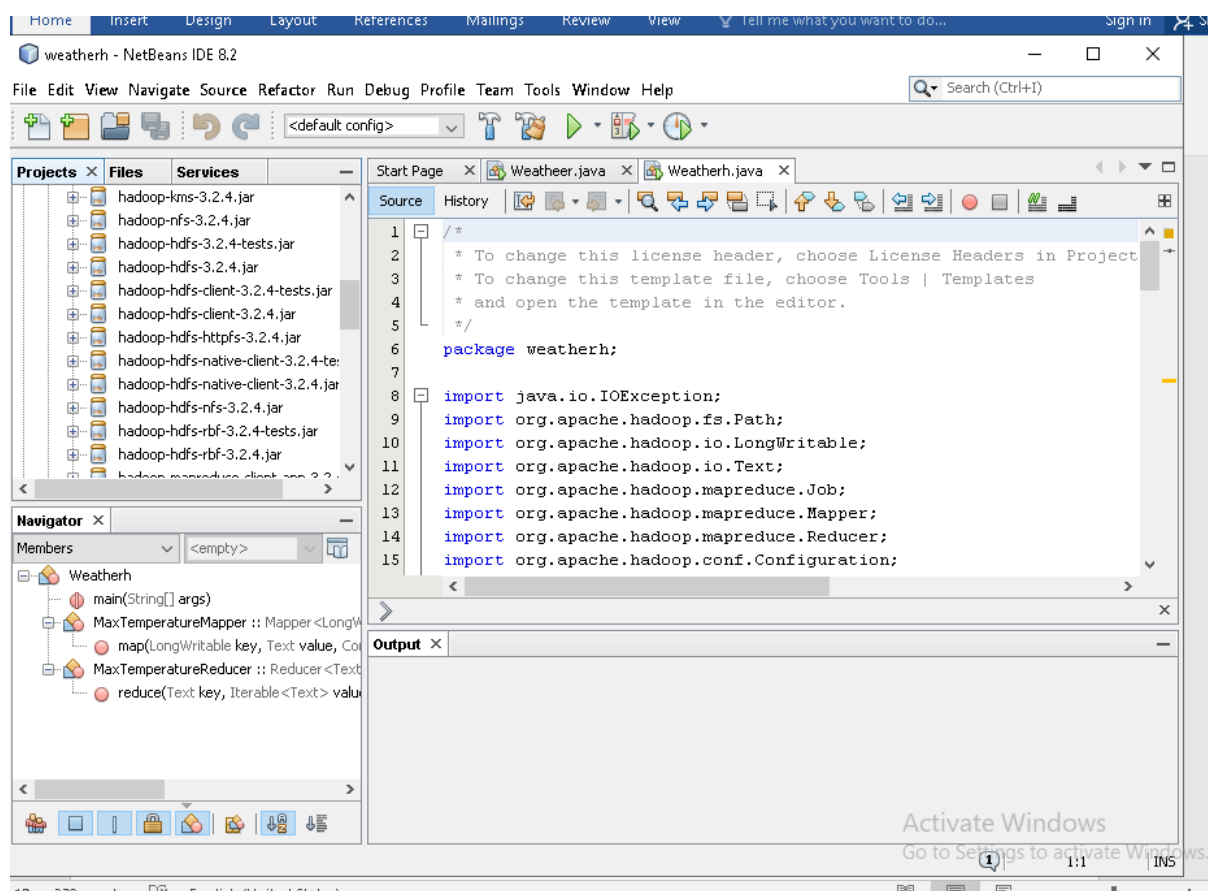
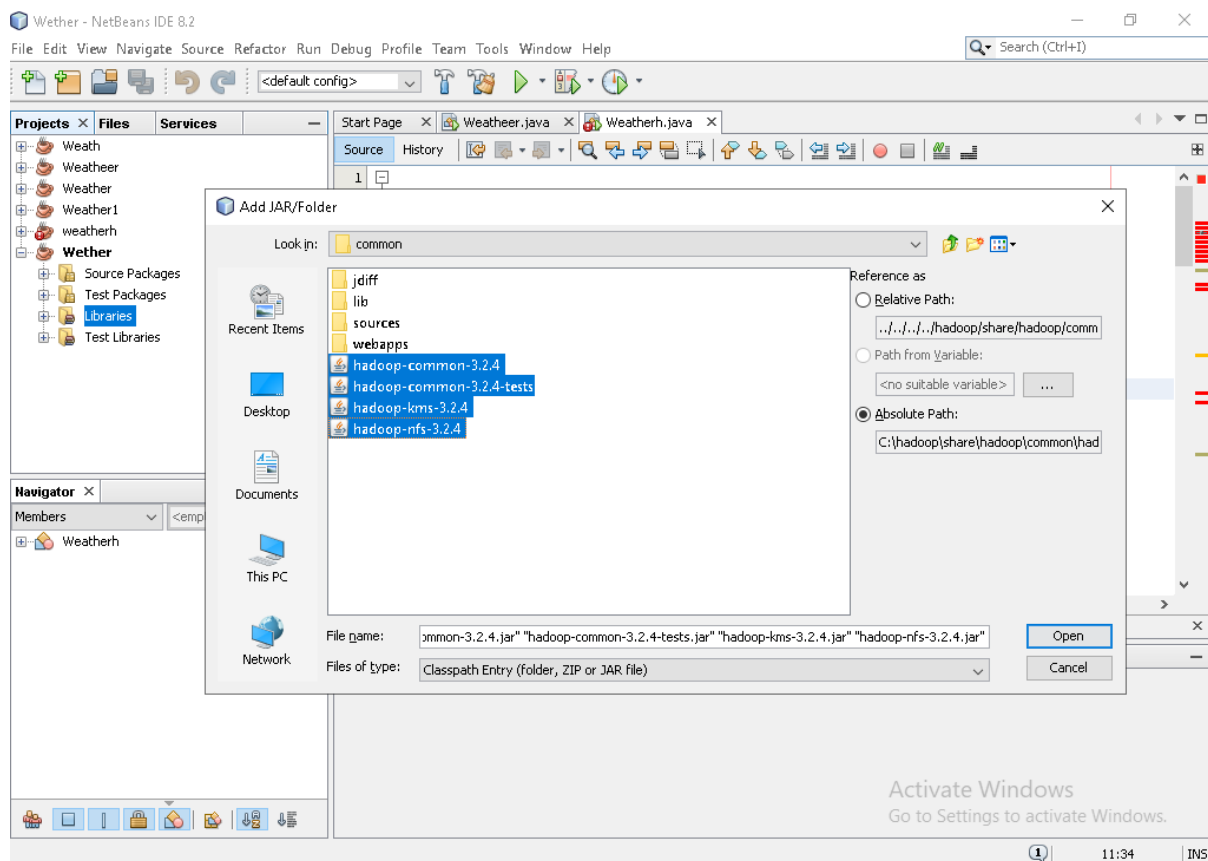
}

}
```

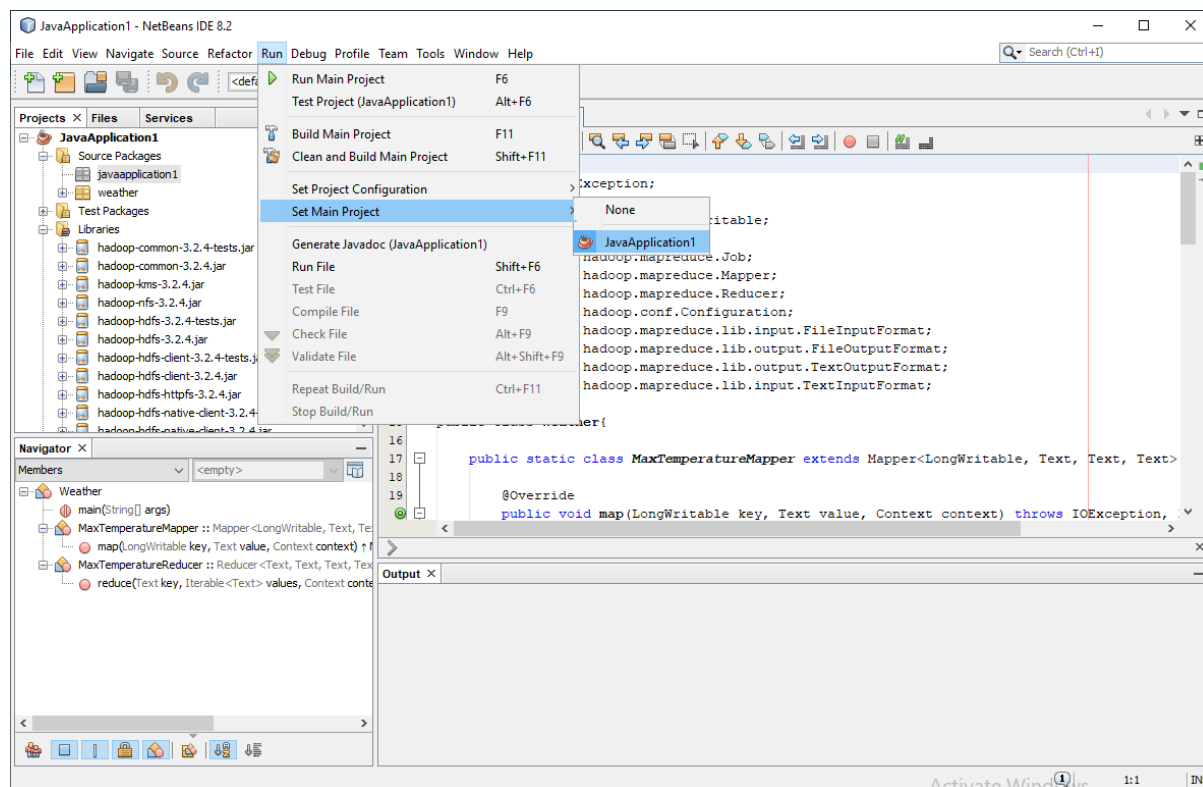
Step 4: Adding the JAR files from Hadoop directory

Go to project panel and select your project name and right click on select ADD JAR/ FOLDER and go Hadoop folder in the directory C:\hadoop\share\hadoop and the import all the JAR files from the folders named common, hdfs, mapreduce and yarn to run the project





Step 5: Building the main project



Now set our project as the Main Project in the netbeans and then Clean and built the main project now a JAR file of the main project will be created in the directory

C:\Users\2023PECAI432\Documents\NetBeansProjects\JavaApplication1\dist

Copy the file to C directory which is the file to run in cmd to get the output.

Step 6: Formating the Namenode

Open Command Prompt as Administrator

C:\Users\Administrator>hdfs namenode –format

Step 7: Start all Hadoop Features

start-all.cmd

start-dfs.cmd

start-yarn.cmd

Step 8: Input and output file

Download weather dataset from any source

Open command prompt and type the following commands to open weather.txt file in the directory named input_dir

Use this directory for input file C:\hadoop\input_dir\weather.txt

```
hadoop fs -mkdir /input_dir
```

```
hadoop fs -put C:/weatherr.txt /input_dir
```

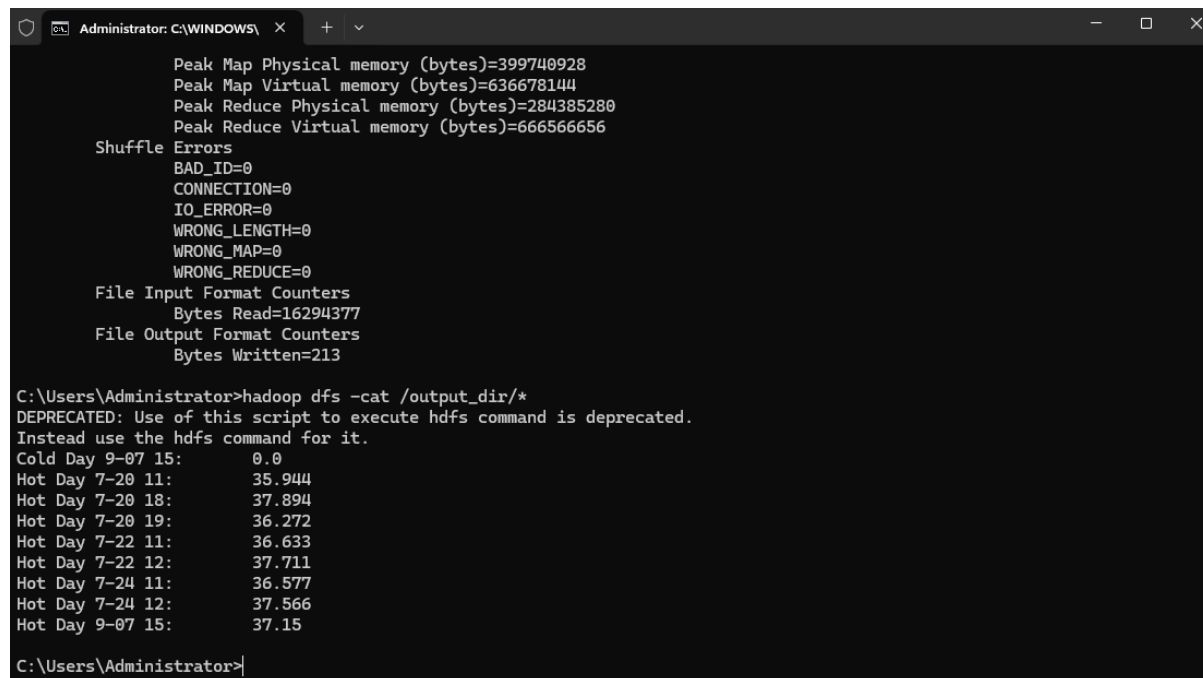
```
Hadoop fs -cat /input_dir/weatherr.txt
```

Use the following commands to show the output in the output_dir directory

```
hadoop jar c:/Weath.jar /input_dir /output_dir
```

```
hadoop jar C:/Weatheer.jar /input_dir /output_dir
```

```
hadoop dfs -cat /output_dir/*
```



The screenshot shows a Windows command prompt window titled "Administrator: C:\WINDOWS\...". The output of the Hadoop commands is as follows:

```
Peak Map Physical memory (bytes)=399740928
Peak Map Virtual memory (bytes)=636678144
Peak Reduce Physical memory (bytes)=284385280
Peak Reduce Virtual memory (bytes)=666566656
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=16294377
File Output Format Counters
  Bytes Written=213

C:\Users\Administrator>hadoop dfs -cat /output_dir/*
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
Cold Day 9-07 15:      0.0
Hot Day 7-20 11:       35.944
Hot Day 7-20 18:       37.894
Hot Day 7-20 19:       36.272
Hot Day 7-22 11:       36.633
Hot Day 7-22 12:       37.711
Hot Day 7-24 11:       36.577
Hot Day 7-24 12:       37.566
Hot Day 9-07 15:       37.15

C:\Users\Administrator>
```

Program :

```
data(mtcars)

# Fit linear regression model
linear_model <- lm(mpg ~ wt, data = mtcars)

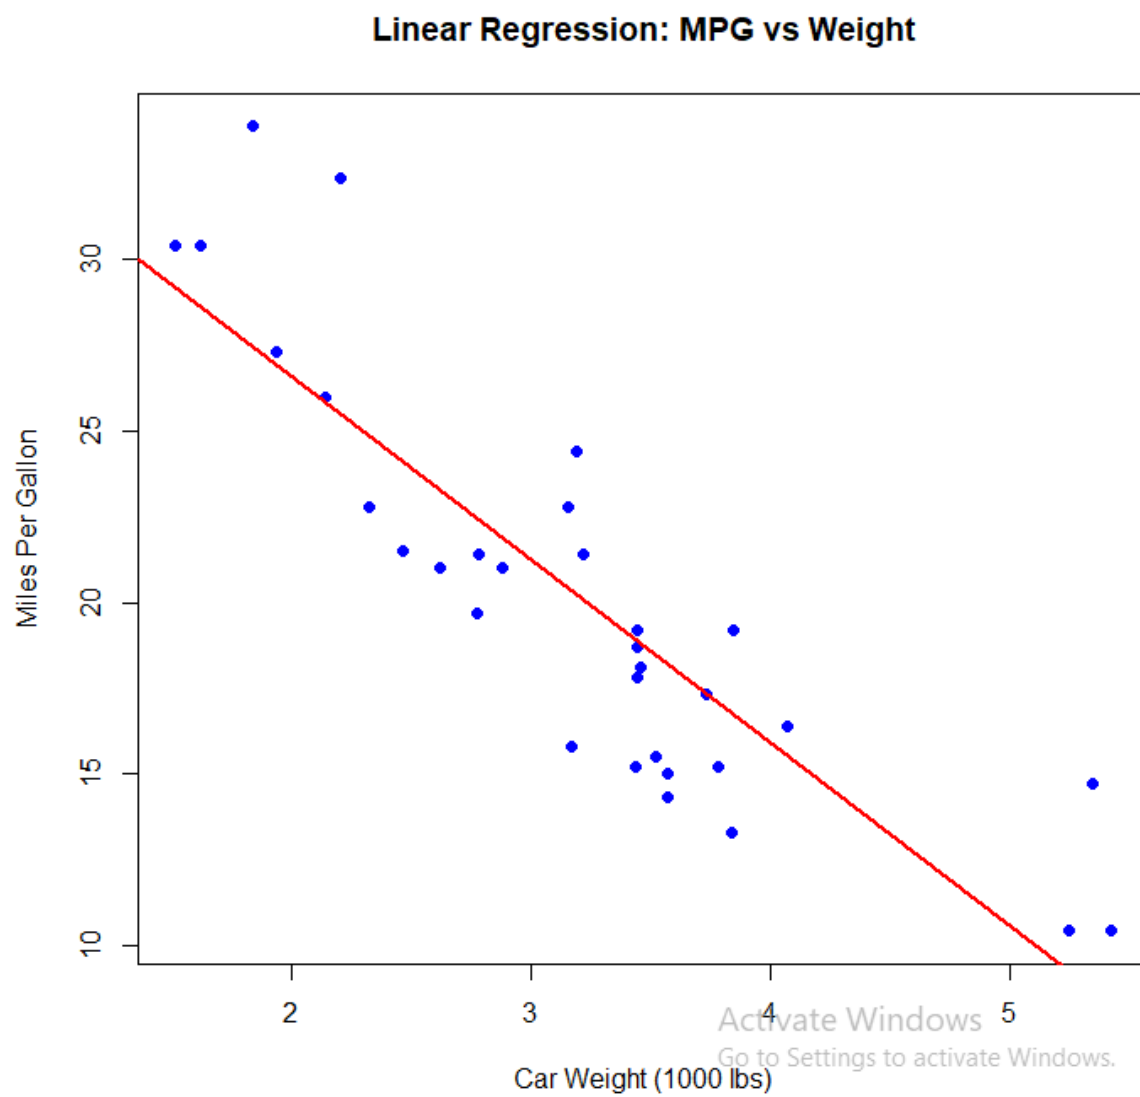
# Summary of the model
summary(linear_model)

# Predict using model
predicted_mpg <- predict(linear_model)

# Plot
plot(mtcars$wt, mtcars$mpg,
     main = "Linear Regression: MPG vs Weight",
     xlab = "Car Weight (1000 lbs)",
     ylab = "Miles Per Gallon",
     pch = 19, col = "blue")
abline(linear_model, col = "red", lwd = 2)

# Calculate R-squared
r_squared <- summary(linear_model)$r.squared
cat("R-squared:", round(r_squared, 3), "\n")
```

Output :



Result :

Program :

```
# Install & load required packages

if (!require("e1071")) install.packages("e1071")

if (!require("ggplot2")) install.packages("ggplot2")


library(e1071)

library(ggplot2)


# Load dataset

data(iris)


# Split data into training (70%) and testing (30%)

set.seed(123)

sample_index <- sample(1:nrow(iris), 0.7 * nrow(iris))

train_data <- iris[sample_index, ]

test_data <- iris[-sample_index, ]


# Train SVM model (linear kernel)

svm_model <- svm(Species ~ ., data = train_data, kernel = "linear")


# Predict on test data

svm_pred <- predict(svm_model, test_data)


# Confusion Matrix

cat("SVM Confusion Matrix:\n")

print(table(Predicted = svm_pred, Actual = test_data$Species))


# 2D Decision Boundary using Sepal.Length & Sepal.Width

svm_model_2d <- svm(Species ~ Sepal.Length + Sepal.Width, data = train_data, kernel = "linear")


xrange <- seq(min(train_data$Sepal.Length), max(train_data$Sepal.Length), length.out = 100)
```

```
yrange <- seq(min(train_data$Sepal.Width), max(train_data$Sepal.Width), length.out = 100)
grid <- expand.grid(Sepal.Length = xrange, Sepal.Width = yrange)
grid$Species <- predict(svm_model_2d, newdata = grid)
```

```
# Plot SVM decision boundaries
```

```
ggplot() +
  geom_point(data = grid, aes(Sepal.Length, Sepal.Width, color = Species), alpha = 0.2) +
  geom_point(data = train_data, aes(Sepal.Length, Sepal.Width, color = Species), shape = 19) +
  labs(title = "SVM Decision Boundaries (2D)",
       x = "Sepal Length", y = "Sepal Width") +
  theme_minimal()
```

Output :



Result :

Program:

```
# Install & load required packages

if (!require("ggplot2")) install.packages("ggplot2")
if (!require("factoextra")) install.packages("factoextra")
if (!require("cluster")) install.packages("cluster")


library(ggplot2)
library(factoextra)
library(cluster)


# Load dataset
data(iris)


# Use only numeric columns (remove Species)
iris_data <- iris[, -5]


# Run K-Means with 3 clusters
set.seed(123)
kmeans_result <- kmeans(iris_data, centers = 3, nstart = 25)


# Display cluster centers
cat("K-Means Cluster Centers:\n")
print(kmeans_result$centers)


# Compare clusters with actual species
cat("\nK-Means Cluster Assignments:\n")
print(table(Actual = iris$Species, Cluster = kmeans_result$cluster))


# Visualize K-Means clustering
fviz_cluster(kmeans_result, data = iris_data,
  palette = c("#2E9FDF", "#00AFBB", "#E7B800"),
```

```
geom = "point", ellipse.type = "convex",  
ggtheme = theme_minimal(),  
main = "K-Means Clustering - Iris Dataset")
```


Output :



Result :

Program :

```
if (!require("ggplot2")) install.packages("ggplot2")

library(ggplot2)

data(iris)

ggplot(iris, aes(x = Sepal.Length, y = Sepal.Width, color = Species)) +
  geom_point(size = 3) +
  labs(title = "Sepal Length vs Sepal Width",
       x = "Sepal Length",
       y = "Sepal Width") +
  theme_minimal()

ggplot(iris, aes(x = Species, y = Petal.Length, fill = Species)) +
  geom_boxplot() +
  labs(title = "Boxplot of Petal Length by Species") +
  theme_minimal()

ggplot(iris, aes(x = Sepal.Length, fill = Species)) +
  geom_histogram(binwidth = 0.3, position = "dodge", color = "black") +
  labs(title = "Histogram of Sepal Length",
       x = "Sepal Length",
       y = "Frequency") +
  theme_light()

if (!require("GGally")) install.packages("GGally")

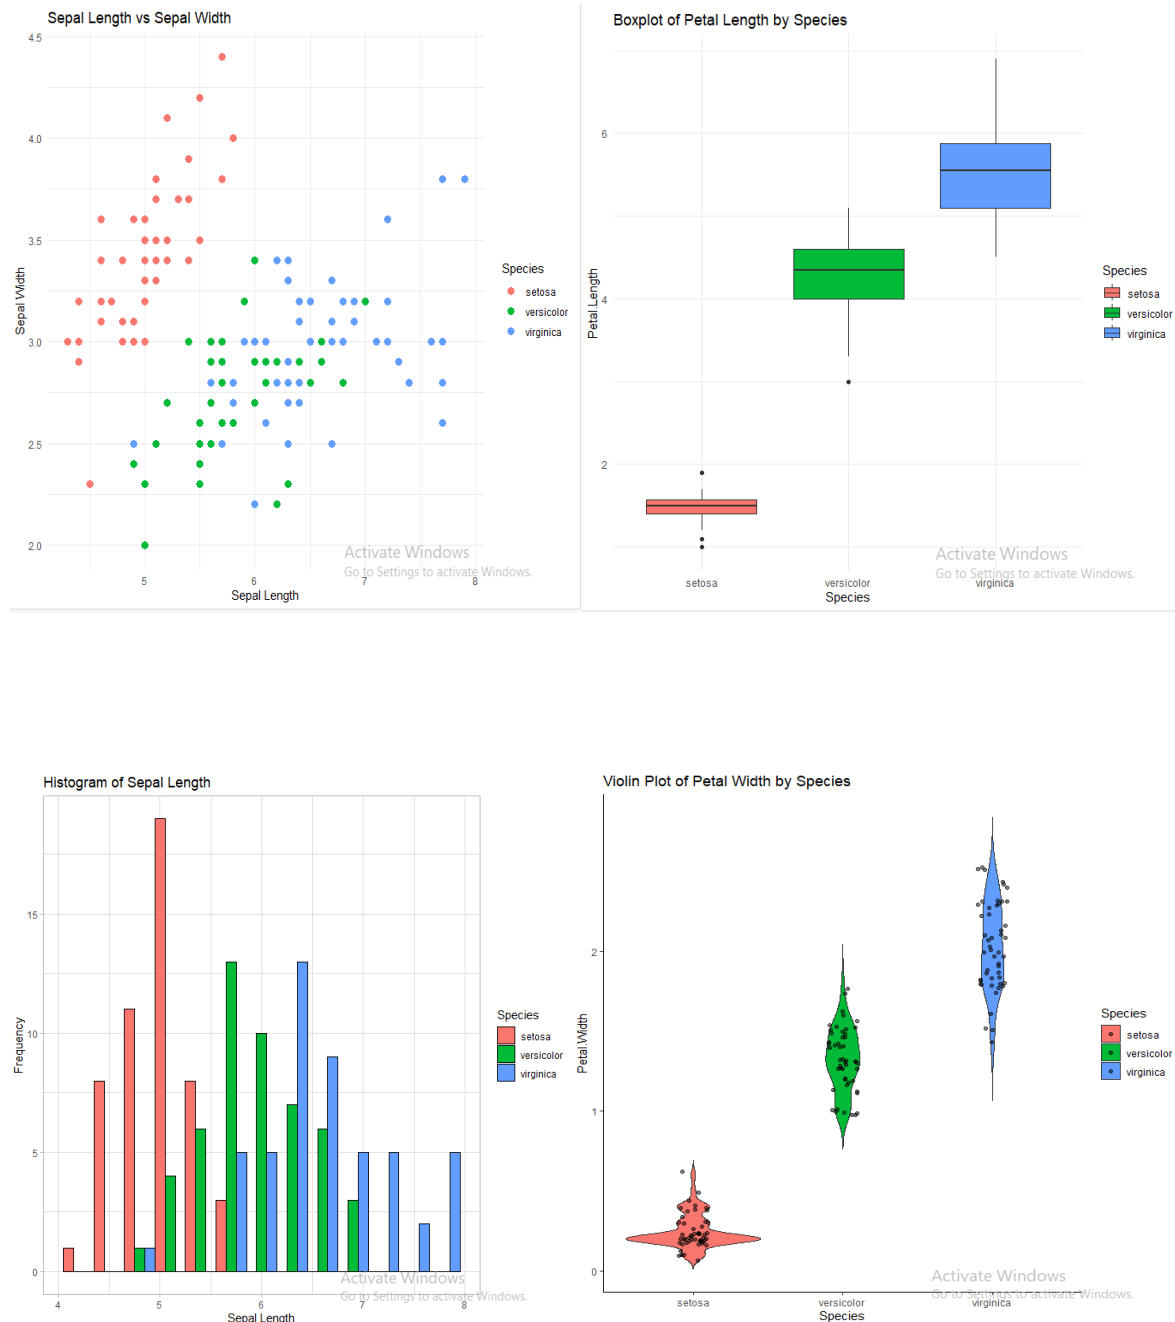
library(GGally)

ggpairs(iris, aes(color = Species),
       title = "Scatterplot Matrix of Iris Features")

ggplot(iris, aes(x = Species, y = Petal.Width, fill = Species)) +
```

```
geom_violin(trim = FALSE) +  
geom_jitter(width = 0.1, alpha = 0.5) +  
labs(title = "Violin Plot of Petal Width by Species") +  
theme_classic()
```

Output :

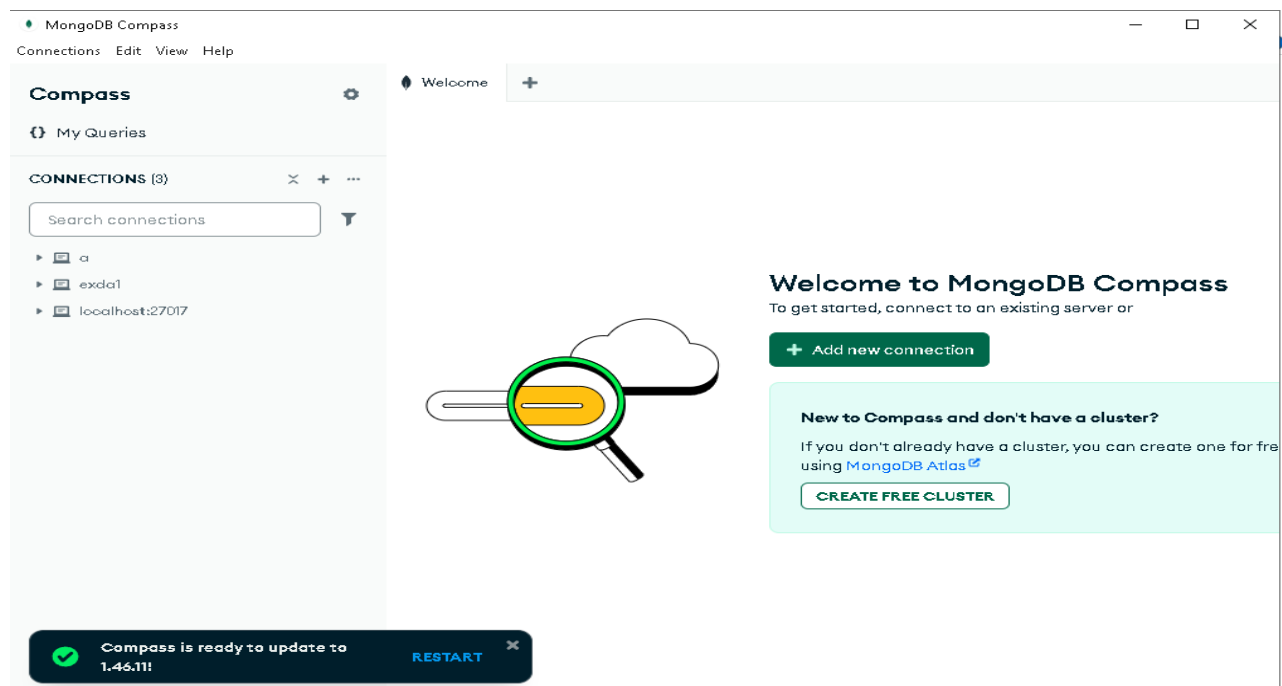


Result :

Program :

Step: 1 <https://www.mongodb.com/try/download/community>

Step2: create mogodb connection , click on Add New connection and sav &e connection



Step 1: Install and load `mongolite` package

```
install.packages("mongolite") # Run only once
```

```
library(mongolite)
```

Step 2: Connect to MongoDB

Replace with your MongoDB URI and database/collection names

```
mongo_conn <- mongo(collection = "students", db = "school", url = "mongodb://localhost")
```

Insert a single document

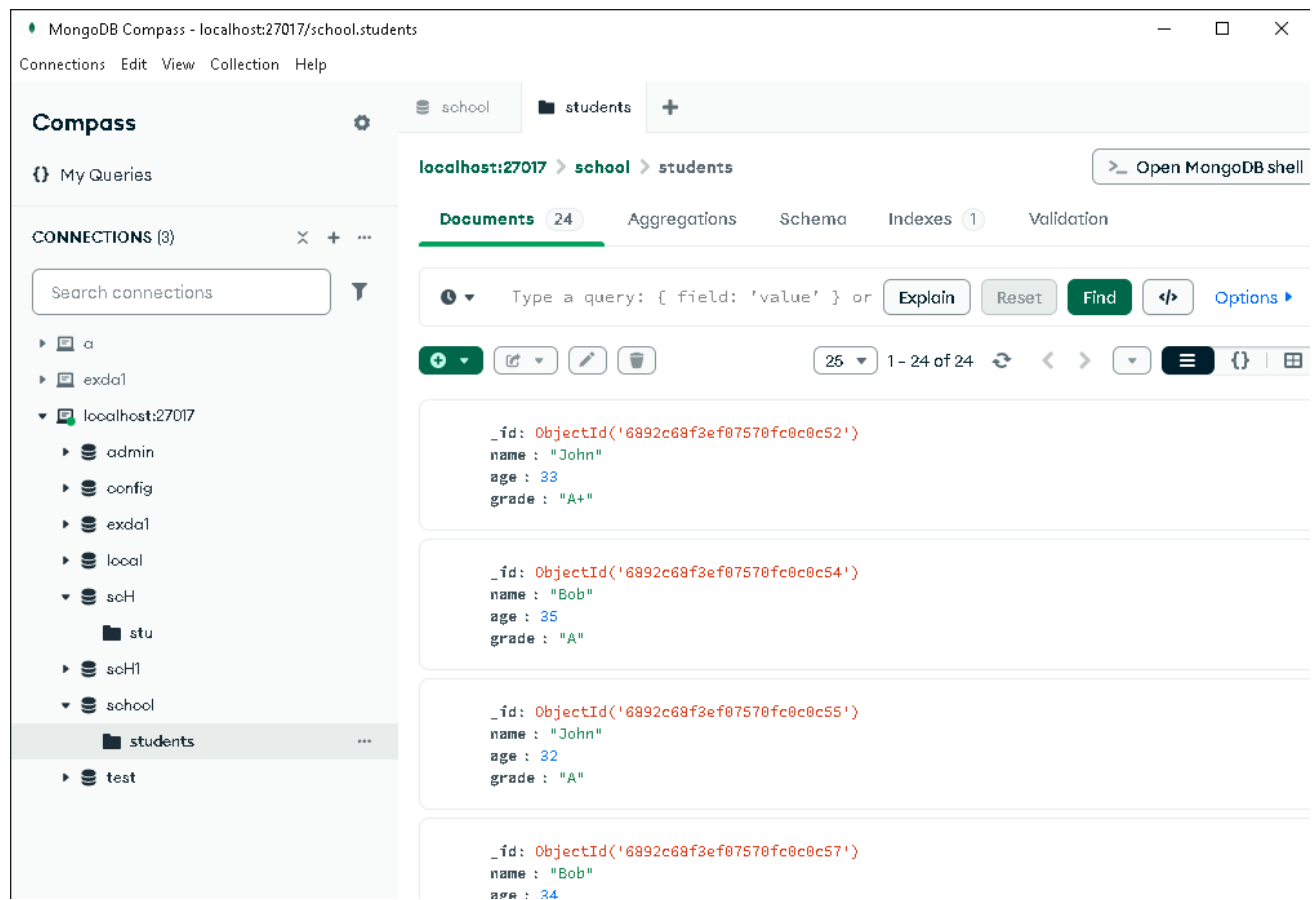
```
mongo_conn$insert('{"name": "John", "age": 21, "grade": "A"}')
```

Insert multiple documents

```
data <- data.frame(
  name = c("Alice", "Bob"),
```

```
age = c(22, 23),  
grade = c("B", "A")  
)  
  
mongo_conn$insert(data)  
  
Find documents  
  
# Find all documents . Retrieve Data  
  
all_docs <- mongo_conn$find()  
  
print(all_docs)  
  
  
# Find with filter  
  
result <- mongo_conn$find({'grade': "A"})  
  
print(result)  
  
Update documents  
  
# Update one document: change grade to "A+" where name is "John"  
  
mongo_conn$update({'name': "John"}, {'$set': {'grade': "A+"}})  
  
  
# Update multiple documents: increase age by 1 for all  
  
mongo_conn$update({'age': {'$exists': true}}, {'$inc': {'age': 1}}, multiple = TRUE)  
  
# Delete one document where name is "Alice"  
  
mongo_conn$remove({'name': "Alice"})  
  
  
# Delete all documents with age > 23  
  
mongo_conn$remove({'age': {'$gt': 23}})  
  
mongo_conn$drop()  
  
mongodb://localhost:27017
```

Output :



MongoDB Compass - localhost:27017/school.students

Connections Edit View Collection Help

Compass

My Queries

CONNECTIONS (3)

Search connections

- alpha
- exda1
- localhost:27017
 - admin
 - config
 - exda1
 - local
 - scH
 - stu
 - scH1
 - school
 - students
 - test

localhost:27017 > school > students

Open MongoDB shell

Documents 24 Aggregations Schema Indexes 1 Validation

Type a query: { field: 'value' } or Explain Reset Find Options

25 1 - 24 of 24

```
{ "_id": "ObjectId('6892c68f3ef07570fc0c52')", "name": "John", "age": 33, "grade": "A+" },
{ "_id": "ObjectId('6892c68f3ef07570fc0c54')", "name": "Bob", "age": 35, "grade": "A" },
{ "_id": "ObjectId('6892c68f3ef07570fc0c55')", "name": "John", "age": 32, "grade": "A" },
{ "_id": "ObjectId('6892c68f3ef07570fc0c57')", "name": "Bob", "age": 34, "grade": "A" }
```

Result :

Program:

```
data(iris)

# Create binary target: 1 if Setosa, 0 otherwise
iris$IsSetosa <- ifelse(iris$Species == "setosa", 1, 0)

# Fit logistic regression model
logistic_model <- glm(IsSetosa ~ Sepal.Length, data = iris, family = "binomial")

# Summary
summary(logistic_model)

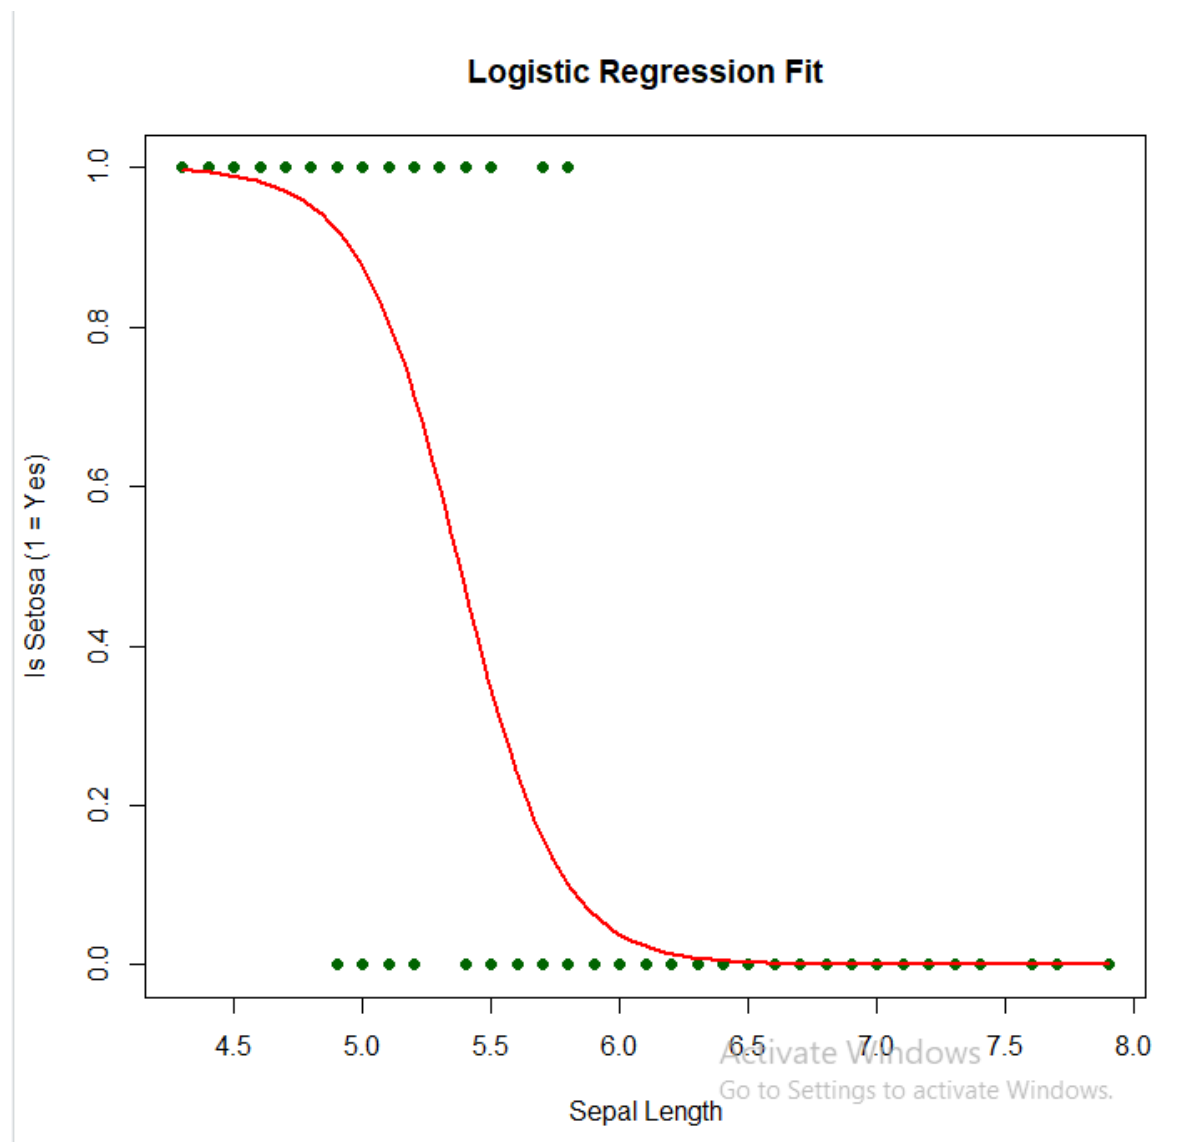
# Predict probabilities
predicted_probs <- predict(logistic_model, type = "response")

# Convert to class labels (0 or 1) with 0.5 threshold
predicted_classes <- ifelse(predicted_probs > 0.5, 1, 0)

# Accuracy
accuracy <- sum(diag(conf_matrix)) / sum(conf_matrix)
cat("Logistic Regression Accuracy:", round(accuracy * 100, 2), "%\n")

# Plot
plot(iris$Sepal.Length, iris$IsSetosa,
     pch = 19, col = "darkgreen",
     xlab = "Sepal Length", ylab = "Is Setosa (1 = Yes)",
     main = "Logistic Regression Fit")
curve(predict(logistic_model, data.frame(Sepal.Length = x), type = "response"),
      add = TRUE, col = "red", lwd = 2)
```


Output :



Result :

Program :

```
# Install & load required packages

if (!require("rpart")) install.packages("rpart")

if (!require("rpart.plot")) install.packages("rpart.plot")


library(rpart)

library(rpart.plot)


# Load dataset

data(iris)


# Split data into training (70%) and testing (30%)

set.seed(123)

sample_index <- sample(1:nrow(iris), 0.7 * nrow(iris))

train_data <- iris[sample_index, ]

test_data <- iris[-sample_index, ]


# Train Decision Tree model

tree_model <- rpart(Species ~ ., data = train_data, method = "class")


# Predict on test data

tree_pred <- predict(tree_model, test_data, type = "class")


# Confusion Matrix

cat("Decision Tree Confusion Matrix:\n")

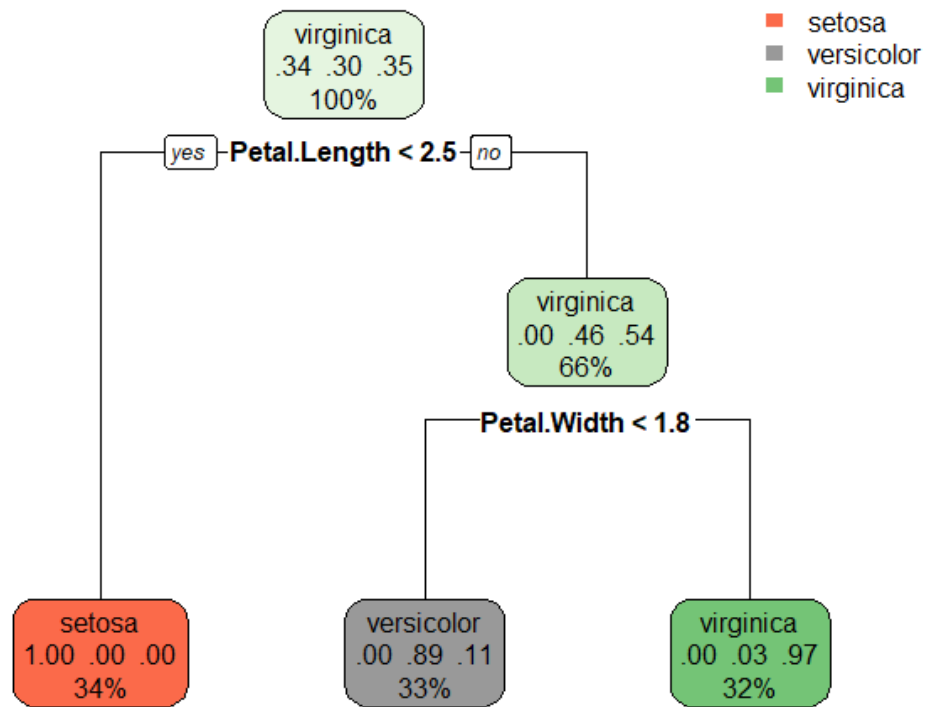
print(table(Predicted = tree_pred, Actual = test_data$Species))


# Plot Decision Tree

rpart.plot(tree_model, main = "Decision Tree for Iris Dataset")
```

Output :

Decision Tree for Iris Dataset



Activate Windows
Go to Settings to activate Windows.

Program :

```
# Hierarchical Clustering on Iris

# Install & load required packages
if (!require("cluster")) install.packages("cluster")

library(cluster)

# Load dataset
data(iris)

# Use only numeric columns (remove Species)
iris_data <- iris[, -5]

# Compute distance matrix
dist_matrix <- dist(iris_data)

# Build hierarchical model using complete linkage
hc_model <- hclust(dist_matrix, method = "complete")


# Plot dendrogram with species labels
plot(hc_model, labels = iris$Species,
     main = "Hierarchical Clustering Dendrogram")

# Cut dendrogram into 3 clusters
hc_clusters <- cutree(hc_model, k = 3)

# Compare clusters with actual species
cat("\nHierarchical Clustering Assignments:\n")
print(table(Actual = iris$Species, Cluster = hc_clusters))
```

Output :

Hierarchical Clustering Dendrogram



Activate Windows

Go to Settings to activate Windows.

dist_matrix
hclust (*, "complete")

Result :